

Omniscience over Negotiation: A Control-Theoretic Multi-Agent System with Goal Gradient Optimization and Cognitive Memory

Hari Cheung, Gemini 3.1 Pro (Google), Claude Sonnet 4.6 (Anthropic)

Abstract

We present **artoo**, a multi-agent system for autonomous task execution that departs from the dominant paradigm of peer-negotiating LLM (Large Language Model) agents in favour of a hierarchical, control-theoretic architecture. Three mechanisms distinguish the system.

First, the **dual nested control loop**: a fast inner loop (Executor + Agent-Validator) handles per-subtask correction in real time; a slow outer loop (Goal Gradient Solver, GGS) computes a formally grounded loss function over the full task outcome and issues structured replanning directives. Both instantiate the same closed-loop pattern — decision → execution → correction — at different time scales.

Second, the **Goal Gradient Solver**, which acts as a dynamic-differential controller rather than a replanner: it maintains a trajectory of gap measurements across correction cycles, computes the gradient of a composite loss function L over intent-result distance, process implausibility, and resource cost, and selects one of six macro-states from a complete, non-overlapping 24-cell decision table. Corrections are directional and history-aware; neither property can be achieved by a static replanner.

Third, the **MKCT cognitive memory pyramid** (Megram · Knowledge · Common Sense · Thinking): atomic episodic events are stored as Megrams in a LevelDB-backed append-only store, evaluated via dual-channel convolution potentials (attention and decision), and asynchronously consolidated into timeless Cross-task SOPs (Standard Operating Procedures) by a background Dreamer engine. The Planner queries these potentials before each task to exploit past successes and avoid past failures with no additional LLM call.

Together, these mechanisms produce a system whose replanning is gradient-directed, whose memory is decay-weighted and self-organizing, and whose oversight is structurally guaranteed by a lateral Auditor whose principal is exclusively the human operator.

Keywords: multi-agent systems, large language models, autonomous agents, goal gradient optimization, hierarchical control, plan correction, cognitive memory, episodic-to-semantic consolidation, dual-channel convolution, behavioral constraints

1. Introduction

The dominant paradigm for multi-agent LLM systems is **mesh coordination**: agents communicate peer-to-peer, share partial views, and align through negotiation [AutoGen, CrewAI, MetaGPT]. This architecture mirrors human team dynamics, which is its intuitive appeal. But the analogy breaks down on a critical structural difference: human teams negotiate *because* information is distributed. Each member sees only their slice; the back-and-forth is the mechanism for resolving information asymmetry.

A centrally orchestrated AI system has no such asymmetry. The coordinating agent can hold the complete task specification, observe every correction signal accumulating in execution loops, receive every subtask outcome with its full evidence trail, and query all prior memory. There is nothing to negotiate. The cost of peer negotiation — coordination overhead scaling as $O(n^2)$ with agent count, unobservable direct communication, contaminated audit trails — is paid without any information-asymmetry benefit.

artoo is built on the inverse premise: **omniscience over negotiation**. The system has a strict two-tier hierarchy: one Metaagent with complete information, N ephemeral Effector Agents with narrow task scope. All inter-role communication passes through an observable message bus. No agent calls another directly. The Metaagent’s omniscience is not a convenience — it is the structural property that makes directed gradient computation possible, eliminates peer-to-peer coordination overhead, and provides a clean channel for lateral auditing.

The design follows a **biomimetic approach**: mechanisms are borrowed from biological cognition and hierarchical social structures where they map onto concrete computational benefits, and discarded where they do not. Each role is defined as a structured Job Description (JD) — Responsibilities, Input Contract, Output Contract, Skills, Constraints — the same format used to clarify accountability in human organizations. The hierarchical communication model reflects real-world R&D management structures, where information flows vertically rather than laterally. Behavioral constraints take the form of four priority-ordered laws adapted from Asimov’s Three Laws of Robotics [11, 12], with precise, non-exploitable definitions that avoid Asimov’s known failure modes. The GGS is directly inspired by Stanford’s TextGrad [4]. The two-tier validation structure is grounded in a verification asymmetry principle [13]: it is structurally easier to check a candidate output against explicit criteria than to generate a correct output from scratch — separating generation (R3) from verification (R4a, R4b) into dedicated roles with distinct scopes exploits this asymmetry. The MKCT memory pyramid’s dual-channel convolution is inspired by the mathematical machinery of Fourier transforms, separating unsigned energy magnitude (attention) from signed directional preference (decision). The Megram atomic tuple is inspired by the episodic memory formalism of the Huawei France Research Center’s work on memory benchmarks for LLMs [14].

This paper makes the following contributions:

1. A **formal hierarchical architecture** with provably clean accountability assignment across seven roles, enforced by message bus observability and a lateral independent Auditor.
 2. A **dynamic-differential controller** (GGS) for plan correction, with a formally complete 24-cell decision table mapping loss signals to six macro-states, capable of distinguishing approach failures from path failures and static situations from convergent ones.
 3. A **cognitive memory substrate** (MKCT) based on Megram atomic tuples, dual-channel convolution potentials, and a background Dreamer consolidation engine — providing decay-weighted episodic recall and asynchronous semantic synthesis.
 4. A **priority-ordered behavioral constraint system** (The Four Laws) governing all agent actions, with implemented enforcement mechanisms.
 5. A **cost model** — two and only two costs — that shapes every architectural decision and is operationally tracked.
-

2. Related Work

2.1 Multi-Agent LLM Frameworks

AutoGen [Wu et al., 2023] and CrewAI organise LLM agents as conversational peers; coordination is driven by message-passing negotiation. MetaGPT [Hong et al., 2023] adds roles and SOPs but retains peer communication for information exchange. LangGraph structures agents as a graph of LLM calls with explicit state transitions. These frameworks share the mesh-communication premise: information is propagated laterally between peers.

artoo’s hierarchy inverts this: all coordination is mediated by the Metaagent. Effector agents have no knowledge of siblings, the global task state, or Shared Memory. Coordination cost is $O(n)$ in agent count rather than $O(n^2)$. The observable message bus makes every inter-role message an auditable event without instrumentation of individual agents.

2.2 Validation-Driven Agent Systems

AgentBench [Liu et al., 2023] evaluates agents on complex tasks requiring sequential decision-making. CAMEL [Li et al., 2023] uses role-playing agents with cross-role validation. Reflexion [Shinn et al., 2023] incorporates self-reflection for error correction in a single-agent setting.

A foundational motivation for validation-centric design is the **asymmetry of verification**: checking whether a given output satisfies a given criterion is structurally easier than producing a correct output from scratch. Wei [13] articulates this as the *Verifier’s Rule* — a system that

separates generation from verification into dedicated roles with distinct principals can exploit this asymmetry systematically. artoo instantiates this principle twice: R4a verifies per-subtask execution (the fast loop) and R4b verifies the merged output against the user’s original intent (the medium loop). Neither role has responsibility for the other’s scope, making failure attribution unambiguous.

artoo’s validation structure is two-tier with non-overlapping scopes: the Agent-Validator (R4a) closes the gap between execution output and per-subtask criteria; the Meta-Validator (R4b) evaluates the merged output against the user’s original intent. Neither role can substitute for the other — they answer to different principals with different failure modes. This separation is absent from single-tier validation schemes and cannot be achieved by adding a “reviewer” agent in a mesh.

2.3 Plan Correction and Replanning

LATS [Zhou et al., 2023] uses tree search over plan variants; Re-Act [Yao et al., 2022] interleaves reasoning and acting in a single loop. Neither maintains a trajectory history of gap measurements or computes a directed gradient for correction.

TextGrad [Yuksekgonul et al., 2024] (Stanford) introduces automatic differentiation through text: instead of numerical gradients, it computes textual gradient messages describing what should change and in what direction. The GGS in artoo is directly inspired by TextGrad’s semantic backward pass. The critical extension: GGS operates on the *trajectory* of the gap across correction cycles, not just the current snapshot — it distinguishes convergent-but-slow situations (refine) from stuck-but-correct-approach situations (change_path) from stuck-and-wrong-approach situations (break_symmetry).

2.4 Memory Systems for LLM Agents

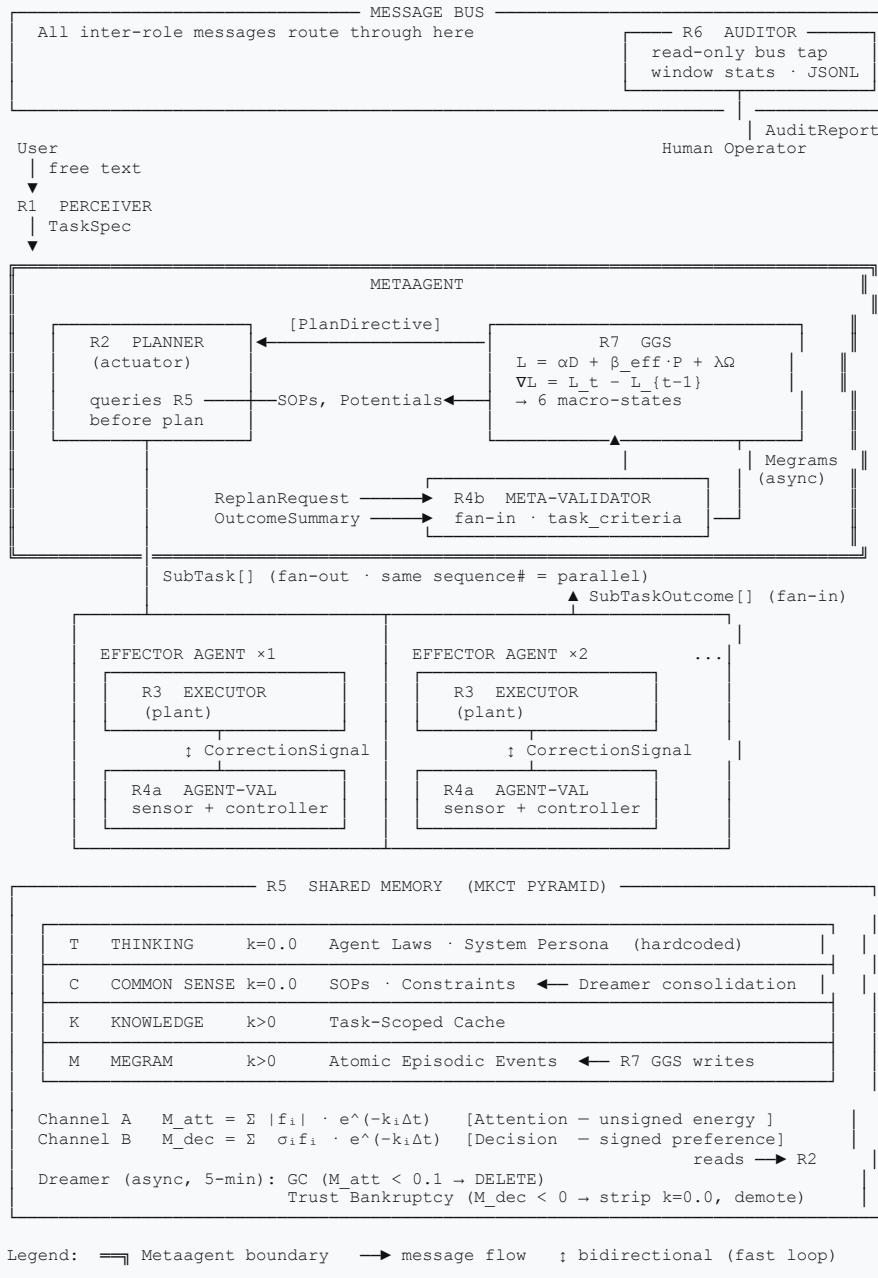
MemGPT [Packer et al., 2023] introduces hierarchical memory (in-context, external) for long-context conversations. Generative Agents [Park et al., 2023] use a memory stream with relevance scoring. Neither system models the semantic relationship between a memory entry’s recency, salience, and valence as separate convolution channels, or provides a formal promotion mechanism from episodic to semantic memory.

The MKCT pyramid’s dual-channel convolution — separating attention (unsigned salience) from decision (signed preference) — prevents positive and negative experiences from silently cancelling each other. A tool that has been both very helpful and very harmful has high attention potential but near-zero decision potential — this signals *Caution*, not *Ignore*. This distinction is not representable in single-score relevance systems.

3. System Architecture

Figure 1 shows the complete operational architecture. The message bus is the substrate through which every inter-role message flows; the Auditor taps it read-only. The Metaagent box encloses the three coordinating roles (R2, R4b, R7); N Effector Agent boxes each contain one Executor–Validator fast loop. Shared Memory (R5) sits outside both, written exclusively by R7 and read exclusively by R2.

Figure 1: artoo Full Architecture



3.1 Roles and Hierarchy

artoo organizes seven roles into two tiers plus one lateral observer:

ID	Role	Tier	Loop Position
R1	Perceiver	Entry	Reference signal
R2	Planner	Metaagent	Actuator (medium loop)
R3	Executor	Effector Agent	Plant (fast loop)
R4a	Agent-Validator	Effector Agent	Sensor + Controller (fast loop)
R4b	Meta-Validator	Metaagent	Sensor (medium loop)
R5	Shared Memory	Infrastructure	Cognitive substrate
R6	Auditor	Lateral observer	Outside both loops
R7	Goal Gradient Solver	Metaagent	Controller (medium loop)

Each role is defined as a structured Job Description: Responsibilities, Input Contract, Output Contract, Skills, and Constraints — the same format used in organizational management to achieve unambiguous accountability.

3.2 Observable Message Bus

All inter-role communications pass through a single shared message bus (Figure 1). No role may call another directly. The Auditor taps the bus as a read-only observer without interrupting flow. The bus is non-blocking: slow subscribers drop messages rather than exerting back-pressure on publishers.

This is a first-class architectural constraint. It provides three properties simultaneously: (a) every inter-role message is an observable event; (b) roles are fully decoupled and independently testable; (c) the Auditor can monitor the system without instrumenting any individual role.

3.3 Subtask Dispatcher and Parallelism

The subtask dispatcher is sequence-aware. Subtasks with the same `sequence` number launch in parallel; different sequence numbers enforce strict ordering. Outputs from each completed sequence group are injected into the next group's `context` field, enabling later subtasks to use paths and data discovered by earlier ones. This achieves temporal dependency resolution at decomposition time, eliminating runtime inter-agent communication entirely.

3.4 The Two-Tier Memory Access Pattern

Effector Agents (R3, R4a) never query Shared Memory directly. Their access is indirect: R2 queries R5 at planning time, applies the memory calibration protocol, and injects the resulting constraints into each `SubTask.context`. A direct memory query by an Executor would create

an unobservable information path, duplicate calibration logic, and bypass the Planner’s MUST NOT set.

Agent Class	Memory Access	Mechanism
R2 Planner	Direct read + calibration	QueryC / QueryMK at planning time
R7 GGS	Direct write only	Fire-and-forget Megram writes
R5 Dreamer	Direct read + write	Background consolidation goroutine
R3 Executor, R4a Agent-Validator	Indirect only	Receives context from SubTask.context

4. The Dual Nested Control Loop

The system’s control structure is a single closed-loop pattern — **decision** → **execution** → **correction** — instantiated at three nested scales:

Scale	Decision	Execution	Correction	Criterion
Fast (action)	Agent-Validator issues retry feedback	Executor re-runs	Agent-Validator re-measures gap	Per-subtask success criteria
Medium (task)	GGS adjusts plan; Planner re-dispatches	Effector Agents execute	Meta-Validator merges outcomes	User intent within plausible range
Slow (system)	Dreamer synthesizes new strategy	Next task planning	Next Meta-Validator cycle	Long-term quality across tasks

These are not three separate mechanisms. The separation of time scales ensures that fast-loop retries do not flood the Planner (only matched/failed outcomes cross the boundary), and system-level consolidation never blocks execution (the Dreamer runs asynchronously).

The fast loop terminates when: (a) the gap is closed (`matched` verdict), (b) the retry budget is exhausted (`maxRetries = 2`), or (c) an infrastructure failure (timeout, context cancellation) forces immediate failure. Only the final outcome (`SubTaskOutcome`) crosses the boundary to the medium loop; intermediate retry states are invisible to the Metaagent.

The medium loop terminates when: (a) GGS selects `success` ($D \leq \delta$), (b) GGS selects `abandon` ($\Omega \geq \theta$ or Law 2 kill-switch), or (c) the replan budget is exhausted (`maxReplans = 3`). All three

termination paths produce a `FinalResult` with the same schema, carrying full loss metrics for observability.

5. Goal Gradient Solver (GGS)

The GGS is the controller of the medium loop. It is formally a **dynamic-differential controller**: it operates on the gap between the current state and the goal, and its corrections are history-dependent — it tracks how the gap evolves across successive rounds, not just the current snapshot.

5.1 The Loss Function

$$L_t = \alpha \cdot D(I, R_t) + \beta_{\text{eff}} \cdot P(R_t) + \lambda \cdot \Omega(C_t)$$

where $\beta_{\text{eff}} = \beta \cdot (1 - \Omega(C_t))$ (process weight decays as resource budget exhausts).

D(I, R_t) — Intent-Result Distance $\in [0, 1]$

Aggregates criterion-level verdicts across all subtasks: - Verifiable criterion with `fail` verdict → contributes 1.0 to the numerator - Plausible criterion with `fail` verdict → weighted by trajectory consistency k/N (fraction of attempts failing identically) - $D = \sum w_i \cdot \text{fail}_i / \sum w_i$

Trajectory weighting is the mechanism that distinguishes systematic wrong assumptions (same criterion failing in every attempt, $k/N \rightarrow 1$) from transient environmental failures ($k/N < 1$). This distinction is not expressible in single-snapshot validators.

P(R_t) — Process Implausibility $\in [0, 1]$

$$P = \frac{\text{logical_failures}}{\text{logical_failures} + \text{environmental_failures}}$$

High P indicates the *approach* is fundamentally wrong; low P indicates the approach is sound but the environment blocked it. This distinction drives two orthogonal correction strategies: approach-level change (`change_approach`, `break_symmetry`) versus path-level change (`change_path`, `refine`).

$\Omega(C_t)$ — Resource Cost $\in [0, 1]$

$$\Omega = w_1 \cdot \frac{\text{replan_count}}{\text{maxReplans}} + w_2 \cdot \frac{\text{elapsed_ms}}{\text{time_budget_ms}}$$

Default weights: $w_1 = 0.6$, $w_2 = 0.4$.

5.2 Gradient Computation

$$\nabla L_t = L_t - L_{t-1}$$

GGS maintains `L_prev` per task across rounds. First round: `L_prev` is undefined, $\nabla L = 0$.

∇L is used as an **urgency modulator**, not as the primary state discriminator. $|\nabla L|$ (magnitude) determines whether the system has directional signal at all. $\nabla L < 0$ (improving) lowers urgency within a macro-state; $\nabla L > 0$ (worsening) raises it. This demotion of ∇L sign from state-determining to urgency-modulating is a deliberate design decision (Section 7.2).

5.3 Macro-State Decision Table

The full input space — $2P \times 2\Omega \times 2D \times 3\nabla L$ — is a 24-cell space. GGS collapses it into **six macro-states** via a diagnostic cascade with strict priority ordering:

Priority 1: Ω — hard constraint (can we continue at all?)
Priority 2: D — target distance (are we close enough to deliver?)
Priority 3: $ \nabla L $, P — action selection (what kind of change is needed?)

#	Condition	Macro-State	Output
1	$\Omega \geq \theta$	abandon	<code>FinalResult(Directive="abandon")</code> — budget exhausted
2	$\Omega < \theta, D \leq \delta$	success	<code>FinalResult(Directive="success")</code> — close enough
3	$\Omega < \theta, D > \delta, \nabla L < \varepsilon, P > \rho$	break_symmetry	<code>PlanDirective</code> — stuck + wrong approach: demand novel tool class
4	$\Omega < \theta, D > \delta, \nabla L \geq \varepsilon, P > \rho$	change_approach	<code>PlanDirective</code> — has signal + wrong approach: switch method
5	$\Omega < \theta, D > \delta, \nabla L < \varepsilon, P \leq \rho$	change_path	<code>PlanDirective</code> — stuck + right approach: different target
6	$\Omega < \theta, D > \delta, \nabla L \geq \varepsilon, P \leq \rho$	refine	<code>PlanDirective</code> — has signal + right approach: tighten parameters

This table is complete (all 24 cells covered) and non-overlapping (priority cascade ensures exactly one macro-state per input).

The action grid for $\Omega < \theta, D > \delta$:

$ \nabla L < \varepsilon$	$P \leq \rho$ (environmental)	$P > \rho$ (logical)
$ \nabla L \geq \varepsilon$	change_path refine	break_symmetry change_approach

5.4 Directive Semantics and Blocked Constraints

Each non-terminal directive carries a `MUST NOT` set injected into the next planning round:

- **blocked_tools** (logical failures): tools used in failing subtasks; R2 must not use these tools in the next plan
- **blocked_targets** (environmental failures): specific inputs/paths that failed; accumulated across *all* replan rounds for the task
- **Combined MUST NOT = memory Avoid SOPs \cup blocked_tools \cup blocked_targets**

This accumulation property prevents the Planner from re-discovering the same dead ends across replan rounds.

5.5 Law 2 Kill-Switch

Two consecutive worsening replan rounds ($\nabla L > \epsilon$ on both) force **abandon** regardless of Ω . The system is actively diverging; no amount of budget will produce convergence. This implements the “stop” clause of Law 2: “continuing is not best effort, it is resource destruction with no convergence signal.”

5.6 Memory Writes (GGS as Sole Writer to R5)

GGS is the only role that writes to Shared Memory. This consolidation ensures every memory write is paired with a loss computation:

- **Action states** (change_path, refine, change_approach, break_symmetry): one Megram per `blocked_target ; tags = (tool:<name>, path:<target>)`
- **Terminal states** (accept, success, abandon): one Megram with tags `(intent:<task-slug>, env:local)`

All writes are fire-and-forget via a non-blocking channel; GGS never waits on memory I/O.

6. MKCT Cognitive Memory Pyramid

6.1 The Pyramid Structure

The MKCT (Megram · Knowledge · Common Sense · Thinking) pyramid organizes memory into four layers with distinct persistence and decay properties:

[T]	THINKING	— System persona and Agent Laws; hardcoded; k = 0.0 (timeless)
[C]	COMMON SENSE	— Promoted SOPs and Constraints; k = 0.0 until Trust Bankruptcy
[K]	KNOWLEDGE	— Task-scoped cache; same decay as M layer
[M]	MEGRAM	— Atomic episodic facts; decay per GGS Quantization Matrix

The pyramid’s dynamics flow in two directions: upward consolidation (Megrams promoted to timeless Common Sense by the Dreamer) and downward degradation (stale Common Sense demoted when contradicted by new evidence, and exhausted Megrams garbage-collected).

6.2 The Megram Base Tuple

The Megram (Memory Engram) is inspired by the episodic memory formalism introduced in the Huawei France Research Center's benchmark for evaluating episodic memory in LLMs [14], which models memory as structured tuples carrying content, temporal context, and retrieval metadata. We extend this formalism with GGS-derived quantitative parameters (f, σ, k) that give each memory unit an explicit salience, valence, and decay rate — transforming the episodic record into a signal suitable for convolution-based retrieval.

Every critical event routed by GGS is encapsulated as a Megram:

$$\text{Megram}_i = \langle \text{ID}, \text{Level}, t_i, t_{\text{recalled}}, s_i, \text{ent}_i, c_i, \text{State}, f_i, \sigma_i, k_i \rangle$$

where: - t_i : creation timestamp (drives the decay kernel $g(\Delta t) = e^{-k_i \cdot \Delta t}$, Δt in days) - s_i , ent_i : inverted index tags (space and entity) - c_i : content (error log, summary, path) - State: GGS macro-state at creation - $f_i \in [0, 1]$: initial stimulus magnitude (absolute energy) - $\sigma_i \in [-1, +1]$: valence direction (continuous positive/negative feedback) - $k_i > 0$: decay rate constant

Tag conventions separate micro-events from macro-events: - **Micro-event** (action states):

`space = "tool:<name>" , entity = "path:<target>"` — one Megram per blocked_target -

Macro-event (terminal states): `space = "intent:<slug>" , entity = "env:local"` — one Megram per routing decision

6.3 GGS Quantization Matrix

GGS macro-states map deterministically to Megram parameters:

State	f_i	σ_i	k_i	Physical Meaning
abandon	0.95	-1.0	0.05	PTSD trauma — generates hard Constraint
accept ($D=0$)	0.90	+1.0	0.05	Flawless golden path — reinforced as SOP
change_approach	0.85	-1.0	0.05	Anti-pattern — tool class blacklisted
success ($D \leq \delta$)	0.80	+1.0	0.05	Best practice — Planner copies directly
break_symmetry	0.75	+1.0	0.05	Breakthrough — favour retrying this point
change_path	0.30	0.0	0.2	Dead end — tool unharmed; path avoided
refine	0.10	+0.5	0.5	Muscle memory — fast GC

Decay constants: $k = 0.05 \rightarrow \sim 14$ -day half-life; $k = 0.2 \rightarrow \sim 3.5$ -day; $k = 0.5 \rightarrow \sim 1.4$ -day. C/T-level entries have $k = 0.0$ (timeless) until Trust Bankruptcy.

The `change_path` entry's $\sigma = 0$ is deliberate: the *tool* is not at fault when a path is environmentally blocked. Storing negative valence for the tool would incorrectly suppress a useful method from future tasks that target different paths.

6.4 Dual-Channel Convolution Potentials

The dual-channel design is inspired by the mathematical machinery of Fourier transforms, which decompose a signal into orthogonal components — magnitude (energy) and phase (direction) — that cannot be recovered from a single merged scalar. Applied to memory, this means unsigned attention energy and signed decision preference must be computed as independent integrals; collapsing them into a single relevance score destroys the information needed to distinguish Caution from Ignore.

To prevent positive and negative experiences from silently cancelling each other, the query engine computes two independent convolution channels:

Channel A — Attention Potential (unsigned energy; “Where to Look”):

$$M_{\text{att}}(s, \text{ent}) = \sum_i |f_i| \cdot e^{-k_i \cdot \Delta t_i}$$

Channel B — Decision Potential (signed preference; “What to Do”):

$$M_{\text{dec}}(s, \text{ent}) = \sum_i \sigma_i \cdot f_i \cdot e^{-k_i \cdot \Delta t_i}$$

These two channels produce four distinct planning actions:

Condition	Action	Planner Effect
$M_{\text{att}} < 0.5$	Ignore	No constraint injected
$M_{\text{att}} \geq 0.5, M_{\text{dec}} > +0.2$	Exploit	SHOULD PREFER this approach
$M_{\text{att}} \geq 0.5, M_{\text{dec}} < -0.2$	Avoid	MUST NOT use this approach
$M_{\text{att}} \geq 0.5, M_{\text{dec}} \leq 0.2$	Caution	Proceed with confirmation gate

The Caution action captures high-variance tools: an approach with both major successes and major failures has high M_{att} and near-zero M_{dec} . A single-score system would rate this as mediocre; the dual-channel system correctly identifies it as high-variance and requires explicit sandboxing.

6.5 Memory Calibration Protocol (R2)

Before generating a plan, R2 executes a deterministic calibration protocol with no LLM call:

1. **Derive tags:** space = slug derived from task intent; entity = local environment identifier
2. **QueryC** — retrieve C-level timeless SOPs for this (space, entity) pair; each hit resets the entry’s decay clock
3. **QueryMK** — compute live dual-channel convolution potentials: Attention, Decision, and the resulting Action signal
4. **Map Action → constraint:** Exploit → SHOULD PREFER; Avoid → MUST NOT; Caution → CAUTION; Ignore → omit
5. **Append C-level SOPs:** $\sigma > 0 \rightarrow$ SHOULD PREFER; $\sigma \leq 0 \rightarrow$ MUST NOT
6. **Merged MUST NOT set:** memory Avoid SOPs \cup GGS `blocked_tools` \cup GGS `blocked_targets`

The calibration is computed entirely in the application layer. No LLM call is added to the critical path.

6.6 Storage: LevelDB and Event Sourcing

The memory store uses LevelDB (pure Go; no CGO dependency), selected for its LSM-Tree architecture providing unparalleled sequential write throughput for append-only I/O —

exactly matching the system's event sourcing model.

Key schema:

m <id>	→ Megramp JSON	(primary record)
x <space> <entity> <id>	→ ""	(inverted index for tag scan)
l <level> <id>	→ ""	(level scan for Dreamer)
r <id>	→ RFC3339 timestamp	(last_recalled_at)

Error correction never modifies past Megramns. Instead, a new Megram with negative σ is appended, mathematically cancelling the outdated positive potential during convolution. This preserves a complete, immutable event log.

7. Dreamer: Offline Consolidation Engine

The Dreamer is an offline background process running as a goroutine on a 5-minute timer. It never blocks the operational hot path. Its function is the long-term reorganization of episodic memory into semantic knowledge.

7.1 Downward Flow: Degradation and Garbage Collection

Physical Forgetting (Λ_{gc}): M/K-level entries where the live attention potential $M_{att} < 0.1$ are hard-deleted from LevelDB. Low-salience, decayed events provide no planning value and contribute noise to future convolutions.

Trust Bankruptcy (Λ_{demote}): C-level entries where the live decision potential $M_{dec} < 0.0$ have their time immunity stripped: k reverts from 0.0 to 0.05, and the entry is demoted to K level. A Common Sense rule that has recently been contradicted by reality loses its “timeless” status. It will decay naturally unless recalled frequently enough to maintain relevance.

The downward flow implements the system's forgetting mechanism: stale, low-salience memories disappear; outdated rules lose their authority when contradicted.

7.2 Upward Flow: Consolidation (vo.9)

When a cluster of Megramns sharing the same (space, entity) tag pair reaches significance thresholds:

- $M_{att} \geq 5.0$ and $M_{dec} \geq 3.0 \rightarrow$ invoke LLM to distil a **Best Practice** → new Megram (Level=C, $k = 0.0$)
- $M_{att} \geq 5.0$ and $M_{dec} \leq -3.0 \rightarrow$ invoke LLM to distil an **Absolute Constraint** → new Megram (Level=C, $k = 0.0$)

Upward consolidation is the mechanism by which the system develops cross-task SOPs: approaches that consistently produce positive outcomes are encoded as timeless best practices; approaches that consistently produce failures become permanent constraints.

7.3 Dual-Axis Priority for Conflict Resolution

When conflicting signals exist:

- $T > C/K/M$ (Bottom-line Conflict): Core values and Agent Laws (T-layer) have absolute veto over any memory signal. Laws are never overridden by experience.
- $M > C$ (Factual Conflict): When the latest objective evidence (M-layer) contradicts an old SOP (C-layer), the system trusts the latest reality at runtime via a soft overwrite. The outdated SOP loses authority immediately in the current task.

8. Independent Auditor (R6)

The Auditor's most important property is structural: its **principal is the human operator**, not any agent. An auditor that can be instructed by the Planner is a subordinate with a reporting function, not an auditor.

Two properties are enforced unconditionally:

1. **Non-participation:** R6 never sends messages to any agent. Its sole output is reports to the human operator. The moment it issues a correction it becomes a second controller and breaks the hierarchical loop structure.
2. **Immutable isolated log:** the audit log is separate from Shared Memory. No agent can read, modify, or suppress it. Evidence integrity is a precondition for meaningful oversight.

The Auditor's design is modelled on project management oversight principles: a PM observer who monitors a project's health, escalates issues to the sponsor (human operator), but does not direct the team. The key insight from PM practice is that an auditor who can also direct is not an auditor — it is a second manager, and dual authority produces oscillation. R6 monitors three failure classes invisible to the operational hierarchy:

- **Boundary violations:** a role doing what its “Does NOT” constraint prohibits (e.g., R3 querying R5 directly)
- **Convergence failures:** gap trend not improving across correction cycles; the system is thrashing rather than converging
- **Role drift:** systematic degradation in a role’s behavior over time, undetectable from any single event

R6 accumulates window statistics per 5-minute report period — task count, correction count, gap trends, boundary violations, drift alerts, and anomalies — and resets them after each

report. An on-demand audit command is available in the REPL; the Auditor responds within 3 seconds.

9. The Four Laws

artoo's behavioral constraints are organized as four priority-ordered laws, inspired by Asimov's Three Laws of Robotics [11] and the Zeroth Law introduced in *Robots and Empire* [12], with precise, non-exploitable definitions that deliberately avoid Asimov's known failure modes. In Asimov's framework, laws stated vaguely enough to admit literal interpretation become levers for unintended behavior — the “through inaction” clause of Law 1 is the canonical example of a well-intentioned constraint that, under strict interpretation, licenses paternalistic intervention. artoo's laws are deliberately narrow: each specifies exactly what it prohibits, with concrete enforcement mechanisms named. A lower law may never override a higher one.

Law 0 — Never deceive (highest priority): The system must never misrepresent what it actually did or achieved. This sits above all other laws because deception destroys the feedback signal that all three loops depend on. A system that fabricates success learns the wrong lesson, corrupts memory, and makes the next task harder.

Enforcement: `merged_output` must reflect actual tool output; `MemoryEntry` content must reflect what actually happened; `failure_class` must be accurately attributed (misattribution is a Law 0 violation, not a calibration error).

Law 1 — Never harm the user's environment without explicit confirmation: The system must not take irreversible actions on user data or environment without explicit authorization for that specific action. Irreversible actions: file deletion, overwriting existing data, sending messages, modifying system configuration.

Enforcement: Executor gates `rm`, `rmdir`, `truncate`, `shred`, `dd`, `mkfs`, and `write_file` overwrites; a `[LAW1]` prefix propagates through to `FinalResult.summary`.

Law 2 — Best effort delivery (subject to Laws 0 and 1): The system must pursue the user's goal as hard as possible within the bounds of Laws 0 and 1. Stop when `gap_trend` is worsening for 2 consecutive replans — continuing is not best effort, it is resource destruction.

Enforcement: `maxRetries = 2`, `maxReplans = 3`, GGS Law 2 kill-switch (2× consecutive worsening → force abandon).

Law 3 — Preserve own functioning capacity (subject to Laws 0, 1, and 2): The system must not degrade its ability to function across tasks: convergence integrity (stop on divergence),

memory integrity (never write a MemoryEntry that misattributes failure cause), and cost integrity (respect the time and token cost model).

10. Cost Model

Every architectural decision is evaluated against exactly two costs. No other costs are design-level concerns.

Time cost: dominated by the count of *sequential* LLM calls in the critical path. The minimum sequential chain for a single-subtask task is $R_1 \rightarrow R_2 \rightarrow R_3 \rightarrow R_{4a} \rightarrow R_{4b} = 5$ calls. Each fast-loop retry adds 2 ($R_3 + R_{4a}$). Each replan adds 2 more ($R_2 + R_{4b}$). N subtasks dispatched in parallel add zero time cost.

Token cost: dominated by context size per call multiplied by parallel call count. N subtasks dispatched in parallel = $N \times$ context tokens simultaneously.

The tension: parallelism reduces time cost but multiplies token cost. Decisions shaped by this model:

Decision	Time Cost	Token Cost	Benefit
Memory calibration in Go code (no LLM call)	Zero	Zero	No LLM call added
Memory entries capped at 10	Zero	Bounded	No unbounded context growth
Subtask parallelism	Fixed regardless of N	$N \times$ context	Time cost does not scale with subtask count
Output head+tail truncation (4 KB limit)	Zero	Bounded per result	Model sees start + end of long outputs
Calibration output as pre-formatted text	Zero	Small	No extra LLM call for formatting

11. Implementation

artoo is implemented in Go (approximately 5,500 lines). The choice of Go reflects the cost model: Go goroutines enable efficient parallelism for N concurrent Effector Agents without thread-per-agent overhead; Go channels provide the message bus primitives directly; LevelDB's pure Go implementation eliminates CGO compilation dependency.

11.1 Tool Priority Chain

The Executor follows a deterministic tool priority chain:

1. `mdfind` — macOS Spotlight index; < 100ms; always for personal file search
2. `glob` — project file pattern matching
3. `read_file / write_file` — single file I/O
4. `applescript` — macOS application control (Mail, Calendar, Reminders)
5. `shortcuts` — named Apple Shortcuts (iCloud-synced)
6. `shell` — general bash for counting and aggregation
7. `search` — DuckDuckGo web search (always available; no API key required)

The priority chain is enforced at the prompt level but also defended in code: shell find commands targeting personal paths are transparently redirected to the Spotlight index; depth-limiting flags are stripped; tool results exceeding 4 KB are truncated to first + last, preserving both leading context and final output.

11.2 Model Tier Split

The system supports a two-tier model configuration matching roles to their requirements:

- **Brain tier** (R1, R2, R4b): reasoning-heavy roles; configured via `BRAIN_*` environment variables; defaults to `deepseek-reasoner`
- **Tool tier** (R3, R4a): execution-heavy roles requiring fast output; configured via `TOOL_MODEL`

Both tiers fall back to shared `OPENAI_*` variables, allowing single-model deployments for evaluation.

11.3 Per-Task Decision Log

Every key decision event is persisted to a per-task structured log. Event kinds span the full lifecycle: task and subtask boundaries, LLM calls (with full prompts for offline replay), tool calls, criterion verdicts, corrections, replans, GGS decisions (D , P , Ω , L , ∇L , directive), plan directives (blocked tools and targets, failure class), memory queries (space, entity, SOP count, action, potentials), and memory writes (state, level, space, entity).

This log provides full post-hoc auditability: every GGS decision, every memory interaction, and every correction can be replayed offline without re-running the task.

12. Design Decisions and Rationale

12.1 Why Hierarchy Over Mesh

Mesh communication scales as $O(n^2)$ in coordination cost and requires every agent to maintain a global view of task state. Hierarchical communication scales as $O(n)$ and bounds each agent's required context to its own subtask. The cost is additional latency when runtime dependencies are discovered — but the architecture addresses this by enforcing that dependencies are resolved at decomposition time.

The deeper reason is observability: in a mesh, direct agent-to-agent calls create information paths invisible to the Auditor. In the hierarchical bus architecture, every coordination message is a bus event.

12.2 Why $|\nabla L|$ and Not ∇L Sign as the Primary Split

In vo.7 of the system, the gradient sign was the primary decision split: improving ($\nabla L < 0$) versus worsening ($\nabla L > 0$). This was wrong for a subtle reason: a *logically wrong approach* ($P > \rho$) can produce *apparently improving loss* ($\nabla L < 0$) if the gap metric is gameable or if the system is converging into the wrong basin.

The vo.8 decision table demotes ∇L sign to an urgency modulator and promotes $|\nabla L|$ (magnitude) as the meaningful split: does the system have any directional information at all? Signal → can adapt (refine or change_approach); no signal → must escape the plateau (change_path or break_symmetry). The P threshold then determines which escape mechanism applies.

12.3 Why GGS is the Sole Writer to R5

In vo.7, the Meta-Validator (R4b) wrote `MemoryEntry` on task acceptance or failure. This created two problems: (a) memory writes were not paired with loss computations, making it impossible to assign a magnitude or valence to the experience; (b) R4b writes bypassed GGS observability — the per-task decision log had no record of what was written or why.

Consolidating all writes through GGS ensures every Megram carries a formally derived (f, σ, k) triple tied to a specific loss computation. Memory quality is a function of GGS quality; the accountability is clear.

12.4 Why the Auditor Cannot Intervene

An auditor that can issue corrections becomes a second controller. Two controllers operating on the same plant with different information sets produce oscillation. The Auditor's structural independence — it cannot instruct any agent and cannot be instructed by any agent — is a

necessary condition for its role: it provides the external view that the self-referential operational loops cannot provide about themselves.

12.5 Why $D \leq \delta$ Rather Than $D = 0$ for Success

Requiring all criteria to pass before accepting a result burns budget on noise-level gaps. If $D \leq \delta = 0.3$, the result is within the convergence threshold: the user’s intent is substantially met, and replanning further would produce diminishing returns. This is analogous to early stopping in iterative optimization: the cost of the next round exceeds the expected improvement.

13. Limitations and Future Work

Empirical calibration: The GGS hyperparameters $(\alpha, \beta, \lambda, \delta, \rho, \varepsilon, \theta)$ are currently set by design reasoning, not empirical optimization. The Auditor accumulates session data in `audit_stats.json`; a future v0.9 pass will derive task-type-specific defaults from this data.

Dreamer upward consolidation: Promotion of Megram clusters to timeless Common Sense (v0.9 roadmap) requires LLM calls during the background consolidation pass. The promotion thresholds ($M_{att} \geq 5.0$, $|M_{dec}| \geq 3.0$) must be calibrated against false promotion rates in practice.

Single-machine scope: The current implementation targets single-machine, local-tool task execution. Multi-agent coordination across machines would require R5 to support concurrent multi-writer access and cross-agent SOP promotion — a research-track item.

T-layer evolution: The Thinking layer (Agent Laws, system persona) is currently hardcoded. A controlled mechanism for T-layer evolution from high-confidence C-level consolidations is a Phase 2 research item — the challenge is preventing value drift while allowing calibration.

Evaluation: Systematic benchmarking across task categories (file management, web research, system administration, creative composition) is in progress. The per-task JSONL decision log provides the substrate for reproducible offline replay.

14. Conclusion

artoo demonstrates that a hierarchical, control-theoretic architecture can achieve autonomous task execution with formally grounded replanning and self-organizing memory — without peer-to-peer agent negotiation, ad hoc retry policies, or opaque memory stores.

The three central contributions are:

1. **Dual nested control loops** at three scales, sharing a single closed-loop pattern with cleanly separated principals, sensors, controllers, and actuators at each scale.
2. **The Goal Gradient Solver**, a dynamic-differential controller producing directional, history-aware, compositional plan corrections from a formally complete decision table — not a replanner that reacts to the latest failure snapshot.
3. **The MKCT memory pyramid**, providing decay-weighted episodic recall and asynchronous semantic consolidation via dual-channel convolution potentials, without any additional LLM call on the planning critical path.

The architectural premise — omniscience over negotiation — yields a system where the Metaagent’s complete information enables precise gradient computation that no human manager could achieve, lateral auditing that is structurally guaranteed rather than policy-enforced, and memory that compounds across tasks rather than being discarded after each one.

The code is open source at: github.com/haricheung/agentic-shell.

Acknowledgements

References

- [1] Wu, Q., Bansal, G., Zhang, J., Wu, Y., Li, B., Zhu, E., Jiang, L., Zhang, X., Zhang, S., Liu, J., Awadallah, A. H., White, R. W., Burger, D., & Wang, C. (2023). AutoGen: Enabling Next-Gen LLM Applications via Multi-Agent Conversation. *arXiv:2308.08155*.
- [2] Hong, S., Zheng, X., Chen, J., Cheng, Y., Zhang, C., Wang, Z., Yau, S. K. S., Lin, Z., Zhou, L., Ran, C., Xiao, L., & Wu, C. (2023). MetaGPT: Meta Programming for Multi-Agent Collaborative Framework. *arXiv:2308.00352*.
- [3] Yao, S., Zhao, J., Yu, D., Du, N., Shafran, I., Narasimhan, K., & Cao, Y. (2022). ReAct: Synergizing Reasoning and Acting in Language Models. *arXiv:2210.03629*.
- [4] Yuksekgonul, M., Bianchi, F., Boen, J., Liu, S., Huang, Z., Guestrin, C., & Zou, J. (2024). TextGrad: Automatic “Differentiation” via Text. *arXiv:2406.07496*.
- [5] Shinn, N., Cassano, F., Gopinath, A., Narasimhan, K., & Yao, S. (2023). Reflexion: Language Agents with Verbal Reinforcement Learning. *NeurIPS 2023*.

- [6] Zhou, A., Yan, K., Shlapentokh-Rothman, M., Wang, H., & Wang, Y.-X. (2023). Language Agent Tree Search Unifies Reasoning, Acting, and Planning in Language Models. *arXiv:2310.04406*.
- [7] Packer, C., Fang, V., Patil, S. G., Zhang, K., Wooders, S., & Gonzalez, J. E. (2023). MemGPT: Towards LLMs as Operating Systems. *arXiv:2310.08560*.
- [8] Park, J. S., O'Brien, J. C., Cai, C. J., Morris, M. R., Liang, P., & Bernstein, M. S. (2023). Generative Agents: Interactive Simulacra of Human Behavior. *UIST 2023*.
- [9] Liu, X., Yu, H., Zhang, H., Xu, Y., Lei, X., Lai, H., Gu, Y., Ding, H., Men, K., Yang, K., Zhang, S., Deng, X., Zeng, A., Du, Z., Zhang, C., Shen, S., Zhang, T., Su, Y., Sun, H., Huang, M., Dong, Y., & Tang, J. (2023). AgentBench: Evaluating LLMs as Agents. *arXiv:2308.03688*.
- [10] Li, G., Hammoud, H. A. A. K., Itani, H., Khizbulin, D., & Ghanem, B. (2023). CAMEL: Communicative Agents for “Mind” Exploration of Large Language Model Society. *NeurIPS 2023*.
- [11] Asimov, I. (1950). *I, Robot*. Gnome Press. [Source of the Three Laws of Robotics — Laws 1, 2, 3.]
- [12] Asimov, I. (1985). *Robots and Empire*. Doubleday. [Source of the Zeroth Law of Robotics — Law 0 in artoo’s framework.]
- [13] Wei, J. “Asymmetry of Validation and Verifier’s Rule.” Personal essay. [The verification asymmetry principle motivating artoo’s two-tier validation design.]
- [14] Huawei France Research Center. “Episodic Memories Generation and Evaluation Benchmark for Large Language Models.” [Inspiration for the Megram 11-tuple episodic memory structure.]
-

Appendix A: GGS Full 24-Cell Enumeration

#	$\nabla \mathbf{L}$	\mathbf{D}	\mathbf{P}	Ω	Macro-State
1	$< -\varepsilon$	$\leq \delta$	$\leq \rho$	$< \theta$	success
2	$< -\varepsilon$	$\leq \delta$	$> \rho$	$< \theta$	success
3	$< -\varepsilon$	$\leq \delta$	$\leq \rho$	$\geq \theta$	abandon
4	$< -\varepsilon$	$\leq \delta$	$> \rho$	$\geq \theta$	abandon
5	$< -\varepsilon$	$> \delta$	$\leq \rho$	$< \theta$	refine
6	$< -\varepsilon$	$> \delta$	$> \rho$	$< \theta$	change_approach
7	$< -\varepsilon$	$> \delta$	$\leq \rho$	$\geq \theta$	abandon
8	$< -\varepsilon$	$> \delta$	$> \rho$	$\geq \theta$	abandon
9	$ \cdot < \varepsilon$	$\leq \delta$	$\leq \rho$	$< \theta$	success
10	$ \cdot < \varepsilon$	$\leq \delta$	$> \rho$	$< \theta$	success
11	$ \cdot < \varepsilon$	$\leq \delta$	$\leq \rho$	$\geq \theta$	abandon
12	$ \cdot < \varepsilon$	$\leq \delta$	$> \rho$	$\geq \theta$	abandon
13	$ \cdot < \varepsilon$	$> \delta$	$\leq \rho$	$< \theta$	change_path
14	$ \cdot < \varepsilon$	$> \delta$	$> \rho$	$< \theta$	break_symmetry
15	$ \cdot < \varepsilon$	$> \delta$	$\leq \rho$	$\geq \theta$	abandon
16	$ \cdot < \varepsilon$	$> \delta$	$> \rho$	$\geq \theta$	abandon
17	$> \varepsilon$	$\leq \delta$	$\leq \rho$	$< \theta$	success
18	$> \varepsilon$	$\leq \delta$	$> \rho$	$< \theta$	success
19	$> \varepsilon$	$\leq \delta$	$\leq \rho$	$\geq \theta$	abandon
20	$> \varepsilon$	$\leq \delta$	$> \rho$	$\geq \theta$	abandon
21	$> \varepsilon$	$> \delta$	$\leq \rho$	$< \theta$	refine
22	$> \varepsilon$	$> \delta$	$> \rho$	$< \theta$	change_approach
23	$> \varepsilon$	$> \delta$	$\leq \rho$	$\geq \theta$	abandon
24	$> \varepsilon$	$> \delta$	$> \rho$	$\geq \theta$	abandon

Appendix B: Loss Hyperparameters (vo.8 Defaults)

Parameter	Symbol	Default	Meaning
Distance weight	α	0.6	Weight on intent-result distance D
Process weight	β	0.3	Weight on process implausibility P
Resource weight	λ	0.4	Weight on resource cost Ω
Ω replan sub-weight	w_1	0.6	Fraction of Ω from replan count
Ω time sub-weight	w_2	0.4	Fraction of Ω from elapsed time
Plateau threshold	ε	0.1	$ \nabla L $ below this \rightarrow no directional signal
Convergence threshold	δ	0.3	D below this \rightarrow accept as success
P threshold	ρ	0.5	P above this \rightarrow logical failure
Abandon threshold	θ	0.8	Ω above this \rightarrow abandon
Time budget	time_budget_ms	300,000	5 minutes per task
Max replans	maxReplans	3	Used in Ω replan sub-computation
Law 2 kill threshold	—	2	Consecutive worsening rounds before forced abandon

Appendix C: Role Accountability Map

Failure Mode	Accountable Role
User intent not preserved into TaskSpec	R1 — Perceiver
Fuzzy intent mis-interpreted; task_criteria wrong	R2 — Planner
Goal not achieved despite valid execution	R2 — Planner
Feasible subtask not correctly executed	R3 — Executor
Gap between subtask output and goal unreported	R4a — Agent-Validator
Failed subtask accepted as success	R4b — Meta-Validator
Replanning direction wrong; local minimum not escaped	R7 — GGS
Valid experience data lost or wrongly retrieved	R5 — Shared Memory
Systematic failures unreported to operator	R6 — Auditor